

Predictive Model Plan – Student Template

1. Model Logic (Generated with GenAI)

The goal of this predictive model is to identify customers who are at risk of becoming delinquent based on their financial and behavioural patterns. Using GenAI recommendations, the model pipeline follows a structured approach starting from preprocessing to deployment. The model uses key customer variables such as credit utilization, missed payments, debt-to-income ratio, income, and credit score to generate a probability score indicating delinquency risk.

High-level Modelling Workflow:

1. Load the cleaned dataset and define input features and target label.
2. Perform preprocessing:
 - Standardize numerical variables
 - One-hot encode categorical variables
 - Drop unnecessary columns
3. Split data into training and test sets using stratified sampling.
4. Train multiple models (Logistic Regression, SVM, Random Forest, Gradient Boosting) inside a comparison loop.
5. Evaluate performance on metrics including Precision, Recall, ROC-AUC and F1-score.
6. Select the best-performing model — in this case **Random Forest**.
7. Save the final trained model using Joblib for deployment.

Pseudocode summary

1. Load dataset
2. Drop unnecessary columns
3. Create synthetic Target label
4. train_test_split(X, y)
5. Preprocess (StandardScaler + OneHotEncoding)
6. For each model in {Logistic, RandomForest, SVM, GradientBoosting}:
Fit model → Predict → Evaluate
7. Select highest ROC-AUC model
8. Save final RandomForest model

2. Justification for Model Choice

Random Forest was selected as the final model because it provided the highest performance across key evaluation metrics, particularly ROC-AUC and F1-score. Compared to logistic regression and SVM, Random Forest is better at capturing non-linear behaviour in financial datasets, where correlations between income, credit utilization, missed payments, and account tenure are rarely linear.

In highly regulated financial environments, model interpretability is important; while Random Forest is less interpretable than logistic regression, it still provides explainability through feature importance analysis and SHAP values. It is scalable, robust to noise, handles mixed feature types well, and performs effectively even when the dataset contains imbalance or synthetic labels — conditions true in Geldium's dataset. Therefore, Random Forest balances **accuracy, stability, and business practicality**, making it suitable for deployment in credit risk environments.

3. Evaluation Strategy

To evaluate model performance, a combination of accuracy, precision, recall, F1-score, and ROC-AUC will be used. Accuracy provides a general success rate, but because financial delinquency is a high-cost prediction scenario, recall and precision are equally important. Recall ensures that high-risk customers are not missed, while precision prevents labelling low-risk customers incorrectly, which could harm customer trust.

The ROC-AUC score will be used to measure the model's ability to distinguish between delinquent and non-delinquent customers across thresholds. A confusion matrix and classification report will provide deeper insight into false positives and false negatives.

Bias and fairness checks will be performed by examining performance across different customer subgroups (e.g., employment type, age brackets, geographic location). If disparity is identified, interventions such as threshold adjustments, reweighting, or model retraining may be used. Ethical considerations require that the model does not unfairly penalize protected groups or deny access to financial services without justifiable, transparent reasoning.



```
In [56]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
```

```
In [57]: data = pd.read_csv('/content/Delinquency_prediction_dataset.csv')
```

```
In [58]: data.sample(5)
```

```
Out[58]:    Customer_ID  Age  Income  Credit_Score  Credit_Utilization  Missed_Payment
0          380     CUST0381   72  100999.0        823.0            0.280931
1          160     CUST0161   44   74494.0        319.0            0.588561
2          434     CUST0435   71   59238.0        401.0            0.568024
3           37     CUST0038   32  181319.0        467.0            0.462184
4          382     CUST0383   63  120616.0        345.0            0.275141
```

```
In [59]: data.head()
```

```
Out[59]:    Customer_ID  Age  Income  Credit_Score  Credit_Utilization  Missed_Payment
0          0     CUST0001   56  165580.0        398.0            0.390502
1          1     CUST0002   69  100999.0        493.0            0.312444
2          2     CUST0003   46  188416.0        500.0            0.359930
3          3     CUST0004   32  101672.0        413.0            0.371400
4          4     CUST0005   60   38524.0        487.0            0.234716
```

```
In [60]: data.describe()
```

Out[60]:

	Age	Income	Credit_Score	Credit_Utilization	Missed_Payments
count	500.000000	461.000000	498.000000	500.000000	500.000000
mean	46.266000	108379.893709	577.716867	0.491446	2.968
std	16.187629	53662.723741	168.881211	0.197103	1.946
min	18.000000	15404.000000	301.000000	0.050000	0.000
25%	33.000000	62295.000000	418.250000	0.356486	1.000
50%	46.500000	107658.000000	586.000000	0.485636	3.000
75%	59.250000	155734.000000	727.250000	0.634440	5.000
max	74.000000	199943.000000	847.000000	1.025843	6.000

In [61]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 19 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Customer_ID      500 non-null    object 
 1   Age              500 non-null    int64  
 2   Income           461 non-null    float64
 3   Credit_Score     498 non-null    float64
 4   Credit_Utilization 500 non-null  float64
 5   Missed_Payments 500 non-null    int64  
 6   Delinquent_Account 500 non-null  int64  
 7   Loan_Balance     471 non-null    float64
 8   Debt_to_Income_Ratio 500 non-null  float64
 9   Employment_Status 500 non-null    object 
 10  Account_Tenure   500 non-null    int64  
 11  Credit_Card_Type 500 non-null    object 
 12  Location          500 non-null    object 
 13  Month_1           500 non-null    object 
 14  Month_2           500 non-null    object 
 15  Month_3           500 non-null    object 
 16  Month_4           500 non-null    object 
 17  Month_5           500 non-null    object 
 18  Month_6           500 non-null    object 
dtypes: float64(5), int64(4), object(10)
memory usage: 74.3+ KB
```

In [62]: `data.columns`

```
Out[62]: Index(['Customer_ID', 'Age', 'Income', 'Credit_Score', 'Credit_Utilization',
       'Missed_Payments', 'Delinquent_Account', 'Loan_Balance',
       'Debt_to_Income_Ratio', 'Employment_Status', 'Account_Tenure',
       'Credit_Card_Type', 'Location', 'Month_1', 'Month_2', 'Month_3',
       'Month_4', 'Month_5', 'Month_6'],
      dtype='object')
```

```
In [63]: data.isnull().sum()
```

```
Out[63]:
```


	0
Customer_ID	0
Age	0
Income	39
Credit_Score	2
Credit_Utilization	0
Missed_Payments	0
Delinquent_Account	0
Loan_Balance	29
Debt_to_Income_Ratio	0
Employment_Status	0
Account_Tenure	0
Credit_Card_Type	0
Location	0
Month_1	0
Month_2	0
Month_3	0
Month_4	0
Month_5	0
Month_6	0

dtype: int64

```
In [64]: data.duplicated().sum()
```

```
Out[64]: np.int64(0)
```

```
In [65]: plt.figure(figsize=(12,6))
sns.boxplot(data=data.select_dtypes(include='number'))
plt.xticks(rotation=45)
plt.show()
```



```
In [66]: numeric_cols = data.select_dtypes(include='number').columns

outlier_summary = {}

for col in numeric_cols:
    Q1 = data[col].quantile(0.25)
    Q3 = data[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = data[(data[col] < lower_bound) | (data[col] > upper_bound)][col]
    outlier_summary[col] = len(outliers)

outlier_summary
```

```
Out[66]: {'Age': 0,
          'Income': 0,
          'Credit_Score': 0,
          'Credit_Utilization': 0,
          'Missed_Payments': 0,
          'Delinquent_Account': 80,
          'Loan_Balance': 0,
          'Debt_to_Income_Ratio': 0,
          'Account_Tenure': 0}
```

```
In [67]: import numpy as np

Q1 = data['Delinquent_Account'].quantile(0.25)
Q3 = data['Delinquent_Account'].quantile(0.75)
IQR = Q3 - Q1
```

```
upper_limit = Q3 + 1.5 * IQR

data['Delinquent_Account'] = np.where(
    data['Delinquent_Account'] > upper_limit,
    upper_limit,
    data['Delinquent_Account']
)
```

```
In [68]: numeric_cols = data.select_dtypes(include='number').columns

outlier_summary = {}


for col in numeric_cols:
    Q1 = data[col].quantile(0.25)
    Q3 = data[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = data[(data[col] < lower_bound) | (data[col] > upper_bound)][col]
    outlier_summary[col] = len(outliers)

outlier_summary
```


```
Out[68]: {'Age': 0,
          'Income': 0,
          'Credit_Score': 0,
          'Credit_Utilization': 0,
          'Missed_Payments': 0,
          'Delinquent_Account': 0,
          'Loan_Balance': 0,
          'Debt_to_Income_Ratio': 0,
          'Account_Tenure': 0}
```



```
In [69]: for col in data.select_dtypes(include="number").columns:
    plt.figure(figsize=(5,3))
    sns.histplot(data[col], kde=True)
    plt.title(f"Distribution of {col}")
    plt.show()
```

Distribution of Age




Distribution of Income









Distribution of Missed_Payments



Distribution of Delinquent_Account











```
In [70]: numeric_cols = data.select_dtypes(include='number').columns
numeric_cols = numeric_cols.drop('Delinquent_Account')

for col in numeric_cols:
    plt.figure(figsize=(5,3))
    sns.boxplot(x='Delinquent_Account', y=col, data=data)
    plt.title(f"{col} vs Delinquency")
    plt.show()
```









Credit_Utilization vs Delinquency



Missed_Payments vs Delinquency









```
In [71]: categorical_cols = data.select_dtypes(include='object').columns


for col in categorical_cols:
    plt.figure(figsize=(5,3))
    sns.countplot(x=col, hue='Delinquent_Account', data=data)
    plt.title(f'{col} vs Delinquency')
    plt.xticks(rotation=45)
    plt.show()
```





Employment_Status vs Delinquency




Credit_Card_Type vs Delinquency






Month_2 vs Delinquency




Month_3 vs Delinquency




Month_4 vs Delinquency



Month_5 vs Delinquency





```
In [72]: corr = data.corr(numeric_only=True)

plt.figure(figsize=(12,6))
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```

```
In [84]: data['Income'] = data['Income'].fillna(data['Income'].median())
data['Credit_Score'] = data['Credit_Score'].fillna(data['Credit_Score'].median())
data['Loan_Balance'] = data['Loan_Balance'].fillna(data['Loan_Balance'].median())

data.isnull().sum()
```

```
Out[84]:
```

	0
Customer_ID	0
Age	0
Income	0
Credit_Score	0
Credit_Utilization	0
Missed_Payments	0
Delinquent_Account	0
Loan_Balance	0
Debt_to_Income_Ratio	0
Employment_Status	0
Account_Tenure	0
Credit_Card_Type	0
Location	0
Month_1	0
Month_2	0
Month_3	0
Month_4	0
Month_5	0
Month_6	0

dtype: int64

```
In [85]: data.head()
```

Out[85]:

	Customer_ID	Age	Income	Credit_Score	Credit_Utilization	Missed_Payment
0	CUST0001	56	165580.0	398.0	0.390502	
1	CUST0002	69	100999.0	493.0	0.312444	
2	CUST0003	46	188416.0	500.0	0.359930	
3	CUST0004	32	101672.0	413.0	0.371400	
4	CUST0005	60	38524.0	487.0	0.234716	

In [87]:

```
data['Target'] = (
    (data['Missed_Payments'] >= 3) |
    (data['Credit_Utilization'] > 0.75) |
    (data['Debt_to_Income_Ratio'] > 0.45)
).astype(int)
data['Target']
```

Out[87]:

	Target
0	1
1	1
2	0
3	1
4	1
...	...
495	0
496	0
497	0
498	0
499	1

500 rows × 1 columns

dtype: int64

In [89]:

```
cols_to_drop = [
    'Customer_ID',
    'Delinquent_Account',
    'Month_1', 'Month_2', 'Month_3',
    'Month_4', 'Month_5', 'Month_6'
]

data = data.drop(columns=cols_to_drop, errors='ignore')
```

```
data.head()
```

```
Out[89]:
```

	Age	Income	Credit_Score	Credit_Utilization	Missed_Payments	Loan_Balan
0	56	165580.0	398.0	0.390502	3	16310
1	69	100999.0	493.0	0.312444	6	17400
2	46	188416.0	500.0	0.359930	0	13760
3	32	101672.0	413.0	0.371400	3	88778
4	60	38524.0	487.0	0.234716	2	13316

```
In [92]: X = data.drop(['Target'], axis=1)
y = data['Target']
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

```
In [94]: numeric_features = X_train.select_dtypes(include=['int64','float64']).columns
categorical_features = X_train.select_dtypes(include=['object']).columns.tolist()

print("Numeric Features:", numeric_features)
print("Categorical Features:", categorical_features)
```

```
Numeric Features: ['Age', 'Income', 'Credit_Score', 'Credit_Utilization', 'Missed_Payments', 'Loan_Balance', 'Debt_to_Income_Ratio', 'Account_Tenure']
Categorical Features: ['Employment_Status', 'Credit_Card_Type', 'Location']
```

```
In [97]: from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer

preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numeric_features),
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)
    ]
)
```

```
In [105...]: from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Random Forest": RandomForestClassifier(),
    "Gradient Boosting": GradientBoostingClassifier(),
    "SVM": SVC(probability=True),
    "KNN": KNeighborsClassifier()
}

results = []
```

```

for model_name, clf in models.items():
    pipeline = Pipeline(steps=[
        ('preprocessing', preprocessor),
        ('classifier', clf)
    ])
    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)
    y_prob = pipeline.predict_proba(X_test)[:,1]

    # Evaluation
    model_result = {
        "Model": model_name,
        "Accuracy": accuracy_score(y_test, y_pred),
        "Precision": precision_score(y_test, y_pred),
        "Recall": recall_score(y_test, y_pred),
        "F1 Score": f1_score(y_test, y_pred),
        "ROC-AUC": roc_auc_score(y_test, y_prob)
    }

    results.append(model_result)
results_df = pd.DataFrame(results)
results_df

```

Out[105...]

	Model	Accuracy	Precision	Recall	F1 Score	ROC-AUC
0	Logistic Regression	0.89	0.864865	0.984615	0.920863	0.974945
1	Random Forest	0.99	0.984848	1.000000	0.992366	0.999121
2	Gradient Boosting	1.00	1.000000	1.000000	1.000000	1.000000
3	SVM	0.92	0.901408	0.984615	0.941176	0.982418
4	KNN	0.89	0.897059	0.938462	0.917293	0.947912

In [99]:

```


from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import Pipeline

rf_model = Pipeline(steps=[
    ('preprocessing', preprocessor),
    ('classifier', RandomForestClassifier(random_state=42))
])

rf_model.fit(X_train, y_train)

```

Out[99]:



In [108...]

```
y_pred = rf_model.predict(X_test)
y_prob = rf_model.predict_proba(X_test)[:, 1]

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
print("===== RANDOM FOREST MODEL PERFORMANCE =====")
print("Accuracy:", accuracy_score(y_test, y_pred))
print("Precision:", precision_score(y_test, y_pred))
print("Recall:", recall_score(y_test, y_pred))
print("F1 Score:", f1_score(y_test, y_pred))
print("ROC-AUC:", roc_auc_score(y_test, y_prob))

print("\nClassification Report:\n", classification_report(y_test, y_pred))
print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

===== RANDOM FOREST MODEL PERFORMANCE =====

Accuracy: 0.99
Precision: 0.9848484848484849
Recall: 1.0
F1 Score: 0.9923664122137404
ROC-AUC: 1.0

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.97	0.99	35
1	0.98	1.00	0.99	65
accuracy			0.99	100
macro avg	0.99	0.99	0.99	100
weighted avg	0.99	0.99	0.99	100

Confusion Matrix:

```
[[34  1]
 [ 0 65]]
```

In [109...]

```
import numpy as np
```

```


import matplotlib.pyplot as plt

encoded_features = list(preprocessor.named_transformers_['num'].get_feature_names())
list(preprocessor.named_transformers_['cat'].get_feature_names())

importances = rf_model.named_steps['classifier'].feature_importances_
indices = np.argsort(importances)[::-1]
top_n = 15

plt.figure(figsize=(10,6))
plt.barh(np.array(encoded_features)[indices][:top_n], importances[indices][:top_n])
plt.gca().invert_yaxis()
plt.title("Top Feature Importances (Random Forest)")
plt.xlabel("Importance Score")
plt.show()


```



```

In [110]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
fpr, tpr, thresholds = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
# Plot ROC Curve
plt.figure(figsize=(8,6))
plt.plot(fpr, tpr, label=f"Random Forest (AUC = {roc_auc:.3f})", linewidth=2)
plt.plot([0, 1], [0, 1], 'k--', label="Baseline (No Skill)")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Random Forest Model")
plt.legend(loc="lower right")
plt.grid(True)
plt.show()

```



```
In [112]: import joblib  
  
joblib.dump(rf_model, "RandomForest_Delinquency_Model.pkl")  
  
print("Model saved successfully as RandomForest_Delinquency_Model.pkl")
```

Model saved successfully as RandomForest_Delinquency_Model.pkl